

Corso di Laurea Triennale in Ingegneria e Scienze Informatiche

Multi-platform distributed systems with Aggregate Computing in Kotlin: the case of proximity messaging

Tesi di laurea in:
PROGRAMMAZIONE AD OGGETTI

Relatore
Prof. Pianini Danilo

Candidato
Luca Marchi

Correlatore
Dott.ssa Cortecchia Angela

Abstract

To my family.

Contents

Abstract	iii
1 Introduction	1
1.1 Context	1
1.1.1 Aggregate Programming	2
1.1.2 Collektive	5
1.2 Motivations	5
1.2.1 Problem statement	6
1.2.2 Real world use cases	7
1.2.3 Goal	7
1.3 State of Art	9
1.3.1 Existing systems	9
1.3.2 Aggregate Computing frameworks	12
1.3.3 Lightweight communication protocols	13
2 Architecture and Implementation	17
2.1 Requirements	17
2.1.1 Functional requirements	17
2.1.2 Non-functional requirements	18
2.1.3 Analysis	18
	19
Bibliography	19

CONTENTS

List of Figures

1.1	Our world is increasingly populated with a wide range of computing devices, embedded in our environment and with many opportunities for local and even location-independent interactions on fixed network infrastructures [BPV15].	2
1.2	Layered structure of Aggregate Programming, illustrating the abstraction at each level [BPV15].	4
1.3	Real-world use case of a decentralized proximity-based messaging application built on Aggregate Computing principles. In this scenario, sensors located near the fire detect the event and propagate an alert message to neighboring nodes within the network [SBEK16].	8
1.4	Diagram illustrating how Briar supports private messaging, public forums, and blogs through multiple transport methods such as Bluetooth, Wi-Fi, and the Internet. [RSG ⁺ 11].	10
1.5	FireChat's mesh networking architecture, where messages are relayed through nearby devices using WiFi and Bluetooth, enabling communication without centralized infrastructure.	11
1.6	High-level architecture of the ScaFi framework, illustrating its components [CVAP22].	13
1.7	Protelis architecture.	14
1.8	MQTT architecture, illustrating the publish-subscribe model with clients and a central broker.	15

LIST OF FIGURES

List of Listings

listings/HelloWorld.java	18
------------------------------------	----

LIST OF LISTINGS

Chapter 1

Introduction

1.1 Context

In a world where Internet of Things (IoT) devices are becoming increasingly prevalent, the need for efficient and reliable communication systems is paramount (Figure 1.1).

The interactions between neighboring devices play a crucial role in enabling seamless data exchange and coordination, so there is the need to design networks with infrastructures that support scalability, adaptability and reusability [BPV15]. In the past, it was reasonable to use a programming model that focused on the individual computing device, and its relationship with one or more users. However, as systems have grown in scale and complexity with the number of computing devices rising, this method has become inadequate.

Traditional network architectures, rely heavily on centralized infrastructures, making them unsuitable for scenarios such as disaster recovery or interactions with neighboring devices. The computational model of *Aggregate Computing* provides a promising approach to address these challenges by enabling decentralized and self-organizing systems [VBD⁺19].

Building on this concept, this thesis explores how Aggregate Programming can support a proximity and decentralized messaging system in a network.

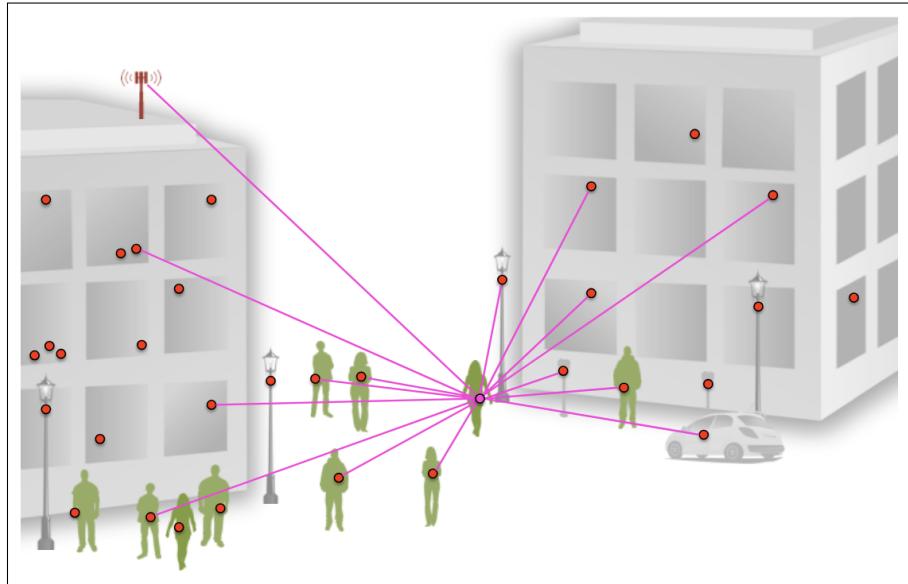


Figure 1.1: Our world is increasingly populated with a wide range of computing devices, embedded in our environment and with many opportunities for local and even location-independent interactions on fixed network infrastructures [BPV15].

1.1.1 Aggregate Programming

Aggregate Programming is a distributed systems paradigm that simplifies programming large networks of devices by focusing on the global, system-level behavior rather than the individual behavior of each device. It shifts the abstraction to view the network as an aggregate, an agglomeration of nodes that collectively exhibit certain behaviors. This model is particularly well-suited for scenarios where devices need to coordinate and collaborate based on local interactions, such as in IoT applications.

As discussed in distributed field programming [VBD⁺19], there are principles to follow when designing large-scale systems:

- The mechanisms ensuring robust coordination should operate transparently in the background, so that programmers don't need to manage them directly.
- The composition of modules and subsystems should be simple and predictable, allowing developers to clearly understand the consequences of combining components.

1.1. CONTEXT

- Large-scale distributed systems often consist of multiple heterogeneous subsystems, each of which may require different coordination strategies depending on the region or context in time.

Aggregate Programming aims to overcome these challenges through three fundamental principles:

1. The computational target is conceived as a region of the environment, with the underlying device-level details abstracted away, potentially even modeled as a continuous spatial domain.
2. The program logic is expressed as the manipulation of data structures that extend spatially and temporally across that region.
3. These computations are executed locally by individual devices within the region, which coordinate through resilient mechanisms and proximity-based interactions.

The foundation of this paradigm lies in *field calculus* [BPV15], a core set of constructs modeling device behavior and interaction. These essential features are captured in a tiny universal language suitable for mathematical analysis. The concept of **field** plays a central role in this context and originates from physics, where it is defined as a function mapping every point in a space–time domain to a scalar value.

According to [VBD⁺19], the *field value* ϕ is a function

$$\phi : D \rightarrow L$$

that maps each device δ in the domain D to a local value ℓ in the range L . This value can change in time with a *field evolution* that maps each point in time to a field value. Figure 1.2 shows how Aggregate Programming abstracts the complexity of the underlying distributed network and its coordination challenges through a sequence of hierarchical abstraction layers.

The messaging system developed in this thesis was implemented using *Collective* [Cor24], an open-source framework designed to simplify the development of

1.1. CONTEXT

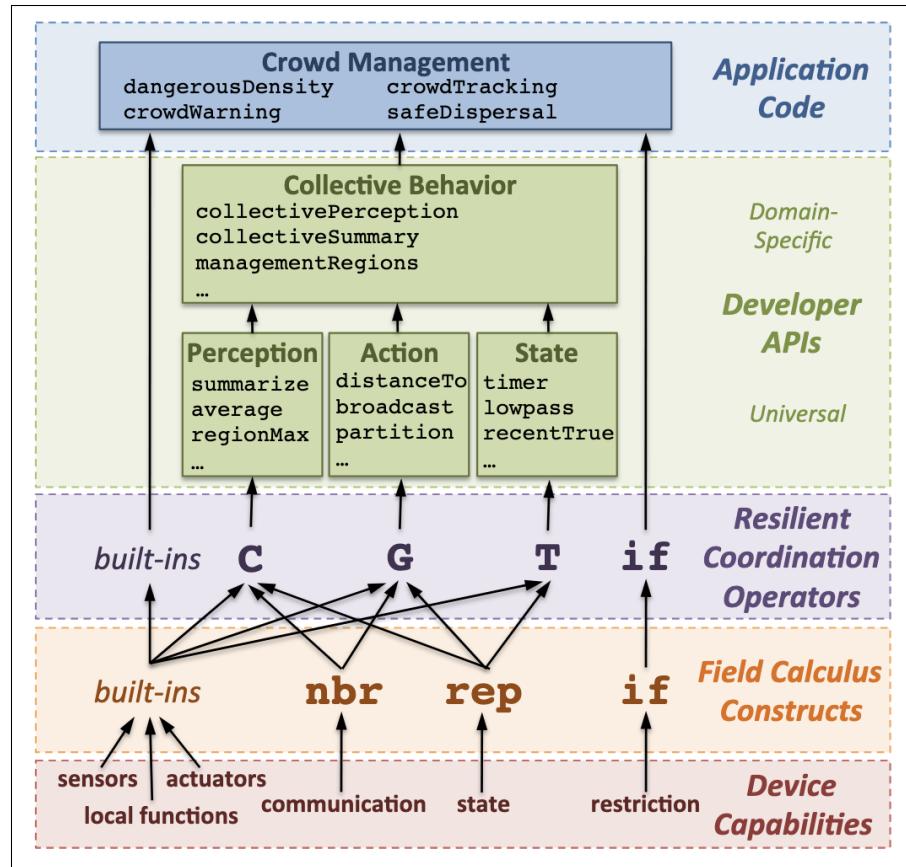


Figure 1.2: Layered structure of Aggregate Programming, illustrating the abstraction at each level [BPV15].

distributed systems through the principles of Aggregate Programming. Collektive provides developer-facing APIs (as in Figure 1.2) used to define and execute aggregate applications.

1.1.2 Collektive

Collektive is a framework with the purpose of simplifying the definition of Aggregate Computing systems. To achieve this, Collektive adopts the Field Calculus model, providing an intuitive way to define aggregate behavior without low-level coding. Built with Kotlin Multiplatform (KMP), it can be executed seamlessly on different targets. Moreover, a compiler plugin was introduced to support function alignment, automatically annotating aligned functions to ensure consistent coordination among nodes [Cor24].

The main features of this framework can be summarized as follows [Cor25]:

- **High-Level Abstraction:** Collektive streamlines the development of distributed behaviors by enabling developers to define coordination logic succinctly, without explicitly handling individual device states.
- **Declarative Programming:** The framework adopts a declarative paradigm, allowing developers to specify the desired system outcomes, while the underlying runtime manages the execution details.
- **Dynamic Adaptation:** Applications developed with Collektive can automatically adapt to network variations and external inputs, maintaining robustness in dynamic and unpredictable environments.

It provides a minimal Domain Specific Language (DSL) where the developer can specify the collective behaviour of a network of devices and the devices can directly communicate with each other and execute the same program.

1.2 Motivations

In this section motivations for implementing a decentralized proximity-based messaging system will be illustrated.

In recent years, the proliferation of IoT and mobile devices has increased the need for local communication systems capable of operating even without stable network infrastructure. Centralized architectures, although efficient in connected environments, fail to guarantee reliability and resilience in dynamic or disconnected scenarios. As devices are becoming pervasive and ubiquitous, new models are required to enable coordination and information exchange without relying on fixed infrastructures. This need motivates the exploration of decentralized and distributed approaches, which can operate robustly in dynamic and connectivity-limited environments.

1.2.1 Problem statement

Traditional messaging systems and architectures often rely heavily on centralized servers and infrastructures, which are unsuitable and vulnerable for scenarios where connectivity is not guaranteed. These situations include **disaster recovery**, the process for restoring an organization's network infrastructure and operations after a disruptive event, such as a natural disaster, cyberattack, or system failure.

In such environments there are several vulnerabilities and challenges to consider:

- **Single point of failure:** Centralized systems are vulnerable to failures of the authority that makes decisions. This can lead to widespread service disruption and outages.
- **Scalability issues:** As the number of devices increases, centralized systems may struggle to manage the growing volume of data and communication, leading to bottlenecks and delayed responses.
- **Limited resilience:** Centralized architectures often lack the redundancy and fault tolerance needed to maintain functionality in the face of network disruptions or device failures.

The networks that need to overcome the issues above can benefit from Aggregate Computing principles, which provides a distributed architecture where the nodes can directly communicate with the neighbors.

1.2.2 Real world use cases

As mentioned in the previous sections, there are several scenarios where a decentralized proximity-based messaging system can be particularly beneficial.

The main use case of such system is in emergency situations. When natural disasters, or infrastructure failures occur, cellular networks and Internet connections often go down, making it difficult for people to communicate and coordinate rescue efforts. Nearby devices can form ad-hoc Bluetooth or Wi-Fi Direct networks, allowing users in the same area to propagate and share messages without relying on centralized servers.

This architecture can also be useful in large IoT networks like sensors systems or drones swarms, operating in remote areas and with the need to communicate without continuous Internet access.

The system's decentralized nature fits perfectly for:

- Exchanging local environmental data (temperature, humidity, motion).
- Maintaining distributed consensus or context awareness among nearby nodes.
- Avoiding single points of failure.

A practical example includes a network of environmental sensors deployed in a forest to monitor conditions and detect wildfires (Figure 1.3), or sensors in a smart agriculture field that share humidity data to coordinate irrigation locally.

1.2.3 Goal

This thesis can be divided into two main goals:

1. **Messaging algorithm:** The first goal is to design and implement a decentralized, proximity-based messaging algorithm grounded in the principles of Aggregate Computing. The network operates in a fully distributed manner, without any central server managing message exchange. Each node acts autonomously, coordinating with its neighbors through local interactions and collectively achieving self-organization, thereby eliminating single points of failure. Communication is proximity-based: devices can only interact with

1.2. MOTIVATIONS

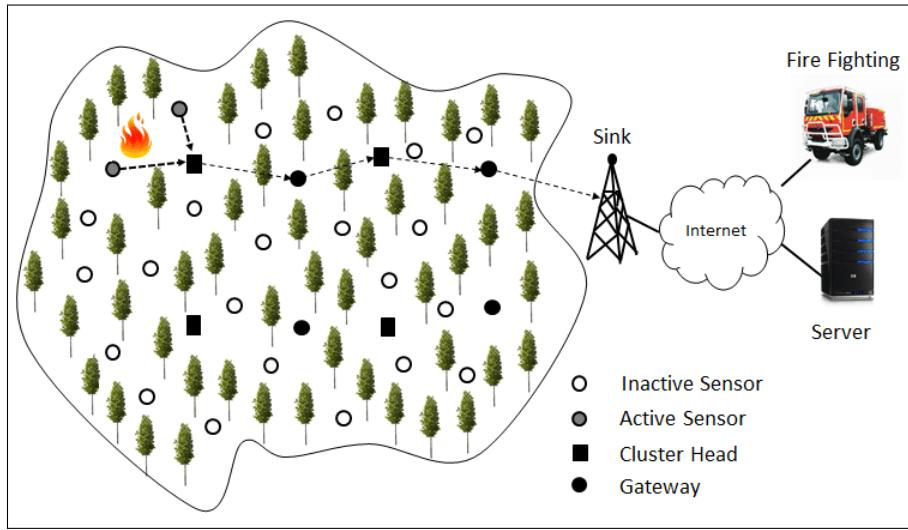


Figure 1.3: Real-world use case of a decentralized proximity-based messaging application built on Aggregate Computing principles. In this scenario, sensors located near the fire detect the event and propagate an alert message to neighboring nodes within the network [SBEK16].

nearby nodes, and messages are propagated progressively across the network, layer by layer. The algorithm must also handle dynamic network topologies, where devices can join or leave the network at any time, ensuring robustness and adaptability. Message sources can define both an expiration time and a maximum propagation distance, enabling controlled diffusion and limiting unnecessary communication overhead. The system is implemented using the *Collektive* framework [Cor24], leveraging KMP to ensure cross-platform compatibility across different targets including JVM, Javascript (browser), Android, iOS, and native versions (for Windows, MacOS, and Linux, both for x86 and ARM CPUs).

2. **Mobile application:** The second goal is to develop a multiplatform mobile application using KMP. This application leverages the messaging algorithm as a core library to send and receive messages from devices in the network. The app must support both Android and iOS platforms providing a simple, user-friendly interface. The User Interface (UI) allows users to compose messages, configure parameters (expiration time and propagation distance),

and view received messages in real-time. Another feature is the anonymization of the users, there is not any registration or login system, the users are completely anonymous. Anybody can use the app without providing any personal information.

1.3 State of Art

In this section, it will be presented the state of the art in the field of decentralized communication systems. The focus will be on the main existing systems, Aggregate Computing frameworks, and lightweight communication protocols.

1.3.1 Existing systems

Several communication technologies and systems have been developed over the years to facilitate decentralized data exchange among devices. The software that has spread the most and gained significant attention includes: *Briar*, *Bridgefy*, *FireChat*.

Briar

Briar [RSG⁺11] is an open-source Peer-to-Peer (P2P) messaging platform designed for secure, resilient communication without centralized servers. It supports multiple transport methods including Bluetooth, Wi-Fi and Tor, with all communication end-to-end encrypted. Briar provide a mailbox system *Briar Mailbox* that lets the users receive encrypted messages from their contacts while Briar is offline. Next time Briar comes online it will automatically fetch the messages from the Mailbox. The primary target audience includes activists, journalists, and civil society members who require secure communication in restricted or infrastructure-limited environments.

A notable feature is the ability to install Briar on an Android device directly from another device that already has it installed, facilitating deployment in disconnected scenarios. Nowadays the application is available only for Android devices, with no official iOS version.

1.3. STATE OF ART

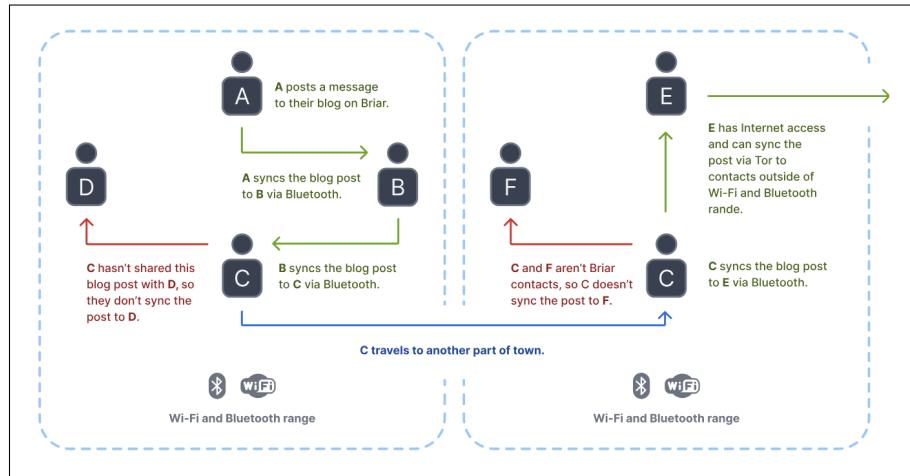


Figure 1.4: Diagram illustrating how Briar supports private messaging, public forums, and blogs through multiple transport methods such as Bluetooth, Wi-Fi, and the Internet. [RSG⁺11].

Bridgefy

Bridgefy [Inc14] is a mobile messaging application that enables offline communication through Bluetooth mesh networking technology.

The platform allows smartphones to connect directly with nearby devices, creating a decentralized mesh network that can relay messages across multiple hops without requiring Internet connectivity. Originally designed for emergency situations and large gatherings where cellular networks become congested, Bridgefy has been deployed in various crisis scenarios including natural disasters and political protests.

Messages are automatically routed through intermediate devices, extending the effective communication range beyond direct Bluetooth connectivity (typically 100 meters). Bridgefy uses encryption protocols to secure message transmission; this capability was not available in the app's initial versions.

The service operates on both Android and iOS devices, with a user-friendly interface designed for rapid adoption during emergency situations. One limitation is the reliance on device density: the mesh network requires sufficient nearby users to establish reliable multi-hop communication paths.

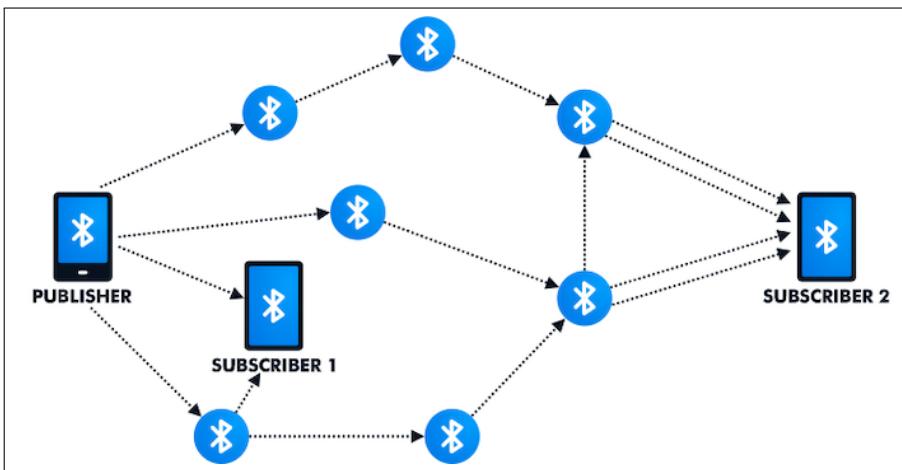


Figure 1.5: FireChat’s mesh networking architecture, where messages are relayed through nearby devices using WiFi and Bluetooth, enabling communication without centralized infrastructure.

FireChat

Firechat was the first popular example of a messaging application that used mesh network topology to communicate between users.

[MTN⁺17] FireChat relays messages from one device to another through each user’s WiFi or Bluetooth radios, thereby creating a “mesh” of phones receiving and relaying messages even when each message is not viewable on every device (Figure 1.5). This method of communication eliminates the need for centralized communication channels such as cellphone towers that every transmission must pass through. While this system is beneficial in overcrowded networks, it eliminates the possibility of controlling or monitoring the network at one point.

Although not originally intended for this purpose, FireChat was widely adopted during civil protests. The application gained popularity in Iraq in 2014, when government restrictions limited Internet access, and later during the 2014 Hong Kong protests (known as the Umbrella Revolution)

Nowadays Firechat is discontinued and no longer available on app stores, the official website is also offline.

Final considerations

Those three illustrated applications have more or less the same purpose of providing a decentralized approach to messaging. *Briar* stands out for its strong emphasis on security and privacy, making it suitable for users in sensitive situations. *Bridgefy* is the most user-friendly and practical, focused on connecting users in emergency scenarios; encryption was included in later versions. *FireChat* was the pioneering first popular example, but now outdated and discontinued. It inspired newer applications.

These applications demonstrate the feasibility of decentralized communication, they rely on predefined routing or broadcasting protocols. None of them incorporate *spatial awareness* or *self-organizing behavior*, features that are naturally modeled through *Aggregate Computing*.

1.3.2 Aggregate Computing frameworks

There are several frameworks and libraries that facilitate the development of Aggregate Computing applications, implementing the principles of field calculus and providing abstractions for defining collective behaviors. Some of the most notable ones include:

- **ScaFi** (Scala Field) is a framework written in Scala for Aggregate Computing [CVAP22]. It provides a DSL and API modules for writing, testing, and running aggregate programs, the high-level architecture of ScaFi is depicted in Figure 1.6. *ScaFi* also includes a simulation environment with a Graphical User Interface (GUI), integrated with the Alchemist simulator [Pia21].

The main applications are in the development of swarm systems, crowd management, wireless sensor network and smart city applications.

ScaFi only supports JVM and web-based applications.

- **Protelis** [PVB15] is a programming language designed to simplify the development of resilient and well-structured networked systems, and can be integrated into various simulation frameworks. It provides a Java-hosted implementation of the field calculus through the `protelis-lang` library, offering

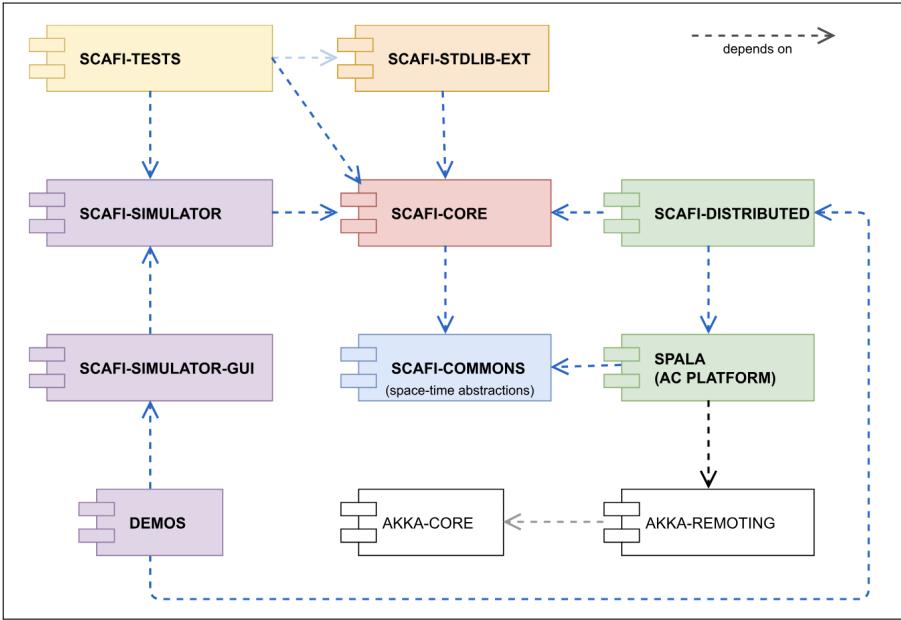


Figure 1.6: High-level architecture of the ScaFi framework, illustrating its components [CVAP22].

a domain-general foundational API for Aggregate Programming [FPBV17]. The Protelis architecture (Figure 1.7) relies on an interpreter that periodically executes pre-parsed aggregate programs, enabling communication among devices and interaction with the environment. However, since it is built on the Xtext framework for domain-specific languages and requires a JVM to run, its deployment is limited to Java-compatible devices, reducing platform heterogeneity.

1.3.3 Lightweight communication protocols

In the context of decentralized and distributed messaging systems, **lightweight communication protocols** are essential for efficient data exchange among devices with limited resources. A lightweight protocol minimizes overhead transmitted on top of the functional data, useful in scenarios where bandwidth, power, or processing capacity is limited. With the proliferation of IoT and mobile devices, these protocols have become crucial in modern networking.

Keeping in mind the distributed nature of the system, we can consider the




Figure 1.7: Protelis architecture.




Figure 1.8: MQTT architecture, illustrating the publish-subscribe model with clients and a central broker.

following technologies.

Bluetooth Low Energy (BLE) formerly marked as *Bluetooth Smart*, is a wireless Personal Area Network (PAN) technology designed for short-range communication with low power consumption. BLE uses the same 2.4 GHz radio frequency as classic Bluetooth but employs a different protocol stack optimized for energy usage. It has no compatibility with classic Bluetooth, meaning that devices supporting only one of the two standards cannot communicate directly.

BLE supports device discovery and connection through low-duty advertising and scanning cycles, enabling the creation of ad-hoc networks for proximity-based communication.

Another similar protocol is **Wi-Fi Direct**, a wireless standard that enables devices to connect directly (using Wi-Fi) to each other without the need for a traditional access point or router. It allows peer-to-peer communication using standard Wi-Fi technology, offering higher data transfer rates and longer range compared to Bluetooth-based solutions.

The table Table 1.1 summarizes the main differences between BLE and Wi-Fi Direct.

On the other hand, **Message Queuing Telemetry Transport (MQTT)** relies on a central broker to manage message distribution. It follows a publish-subscribe model (Figure 1.8) where the centralized server known as Broker is con-

1.3. STATE OF ART

Table 1.1: Comparison between BLE and Wi-Fi Direct.

Feature	BLE	Wi-Fi Direct
Communication range	Typically up to 50 meters (depending on environment)	Up to 200 meters (depending on device and power)
Data rate	Up to 2 Mbps (BLE 5.0)	Up to 250 Mbps
Consumption	Very low, optimized for battery-powered devices	Higher, comparable to standard Wi-Fi
Network topology	Simple connections or small mesh networks	Group-based topology with a dynamically assigned group owner
Connection setup	Fast, lightweight pairing through advertising and scanning	Slower setup, requires negotiation between devices
Use cases	Sensor networks, wearables, proximity-based messaging	File sharing, multimedia streaming, ad-hoc data exchange
Advantages	Energy efficient, widely supported, ideal for IoT	High throughput, longer range, infrastructure-free
Limitations	Limited bandwidth and range	Higher energy consumption, more complex setup

nected to various devices which are Clients [Usm21]. The latter publish messages to specific topics, while subscribers receive only the data relevant to them.

MQTT is widely used in IoT systems due to its lightweight nature and efficient use of bandwidth.

In the current implementation of the proposed application 2, MQTT is employed as the communication layer to handle messages across devices. Although it relies on a central broker and is therefore not fully decentralized, its advantages make it suitable for prototyping and testing the distributed algorithm 1. Future developments may involve replacing MQTT with peer-to-peer technologies such as Bluetooth, to achieve complete decentralization.

Chapter 2

Architecture and Implementation

2.1 Requirements

This section outlines both the functional and non-functional requirements of the messaging system developed in this thesis. It describes what the system is expected to accomplish, its qualities, constraints, and performances.

The requirements refer to both the messaging algorithm itself (see 1) and the mobile application (see 2).

2.1.1 Functional requirements

The functional requirements define the features, behaviors, and operations that the messaging system must support.

Aggregate Computing algorithm

The messaging algorithm represents the Aggregate Computing layer, the core logic of the system.

2.1. REQUIREMENTS

```
1 public class HelloWorld {  
2     public static void main(String[] args) {  
3         // Prints "Hello, World" to the terminal window.  
4         System.out.println("Hello, World");  
5     }  
6 }
```

2.1.2 Non-functional requirements

2.1.3 Analysis

You may also put some code snippet (which is NOT float by default), eg: section 2.1.3.

Bibliography

- [BPV15] Jacob Beal, Danilo Pianini, and Mirko Viroli. Aggregate programming for the internet of things. *Computer*, 48(9):22–30, 2015.
- [Cor24] Angela Cortecchia. *A Kotlin multi-platform implementation of aggregate computing based on XC*. PhD thesis, University of Bologna, 2024.
- [Cor25] Angela Cortecchia. Collektive documentation, 2025.
- [CVAP22] Roberto Casadei, Mirko Viroli, Gianluca Aguzzi, and Danilo Pianini. Scafì: A scala dsl and toolkit for aggregate programming. *SoftwareX*, 20:101248, 2022.
- [FPBV17] Matteo Francia, Danilo Pianini, Jacob Beal, and Mirko Viroli. Towards a foundational api for resilient distributed systems design. In *2017 IEEE 2nd International Workshops on Foundations and Applications of Self* Systems (FAS*W)*, pages 27–32, 2017.
- [Inc14] Bridgefy Inc. Bridgefy: Offline messaging app, 2014.
- [MTN⁺17] Luke P. Morrison, Brian Team, Brian Nguyen, Senthil Kannan, Nathan Ray, and Gregory C. Lewin. Airchat: Ad hoc network monitoring with drones. In *2017 Systems and Information Engineering Design Symposium (SIEDS)*, pages 38–43, 2017.
- [Pia21] Danilo Pianini. Simulation of large scale computational ecosystems with alchemist: A tutorial. In Miguel Matos and Fabíola Greve, editors, *Distributed Applications and Interoperable Systems*, pages 145–161, Cham, 2021. Springer International Publishing.

BIBLIOGRAPHY

- [PVB15] Danilo Pianini, Mirko Viroli, and Jacob Beal. Protelis: practical aggregate programming. In *Proceedings of the 30th Annual ACM Symposium on Applied Computing*, SAC '15, page 1846–1853, New York, NY, USA, 2015. Association for Computing Machinery.
- [RSG⁺11] Michael Rogers, Eleanor Saitta, Torsten Grote, Julian Dehm, Mikolai Güttschow, Nico Alt, and Briar Project. Briar: Secure messaging anywhere, 2011.
- [SBEK16] Massinissa Saoudi, Ahcène Bounceur, Reinhardt Euler, and Tahar Kechadi. Energy-efficient data mining techniques for emergency detection in wireless sensor networks. In *2016 IEEE International Conference on Smart Computing (SMARTCOMP)*, pages 1–6, 07 2016.
- [Usm21] Mohammad Faiz Usmani. Mqtt protocol for the iot - review paper, 05 2021.
- [VBD⁺19] Mirko Viroli, Jacob Beal, Francesco Damiani, Stefano Montagna, Danilo Pianini, Luca Ricci, and Franco Zambonelli. *Aggregate programming: from foundations to applications*. Springer, 2019.